

LECTURE 3

Linear classification

Foundations of Artificial Intelligence

SEMTM0051 · 2026/27

Lecturer: Dr Roberto Colomboni



Today's lecture

Binary classification

Classification error and logistic loss


Gradient descent

More than two classes

Binary classification

A synthetic dataset adapted from Iris:

100 observations $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_{100}, y_{100})$, with two input features and two classes.



$$\mathbf{x}_i = (x_{i1}, x_{i2}), \quad y_i \in \{0,1\}$$

x_1 : first input feature

x_2 : second input feature

Class 0: 50 observations

Class 1: 50 observations

Learn a rule to predict the class of an input.

[Dataset and instructions in the folder: Week 3 - Examples](#)

Measuring classification error

A regression **predictor** $f: \mathbb{R}^d \rightarrow \mathbb{R}$
returns a **number**

Regression

$$\ell_{\text{sq}}(\hat{y}, y) = (\hat{y} - y)^2$$

Empirical squared loss

$$\sum_{i=1}^N \ell_{\text{sq}}(f(\mathbf{x}_i), y_i)$$

Larger numerical errors
receive larger penalties

A **classifier** $h: \mathbb{R}^d \rightarrow \{0,1\}$
returns a class **label**.

Classification

$$\ell_{0,1}(\hat{y}, y) = \begin{cases} 0, & \hat{y} = y \\ 1, & \hat{y} \neq y \end{cases}$$

Number of mistakes / Empirical 0-1 loss

$$\sum_{i=1}^N \ell_{0,1}(h(\mathbf{x}_i), y_i)$$

All mistakes are penalized
the same way

Predicting a probability first

A label gives a **decision**.

A **probability** also records how strongly the model favours that label.

Probabilistic predictor

$$p : \mathbb{R}^d \rightarrow [0,1]$$

Probability of class 1

$$p(\mathbf{x})$$

Probability of class 0

$$1 - p(\mathbf{x})$$

Example: if $p(\mathbf{x}) = 0.51$ we think that there is

- a 0.51 probability that the correct class of input $\mathbf{x}$ is 1
- a 0.49 probability that the correct class of input $\mathbf{x}$ is 0

If we must choose one label, which should we predict?


The plug-in classification rule

Choose the label with the larger probability. At a tie, we choose class 0.

$$h_p(\mathbf{x}) = \begin{cases} 0, & p(\mathbf{x}) \leq 1/2 \\ 1, & p(\mathbf{x}) > 1/2 \end{cases}$$

Heaviside step function

$$\text{step}(s) := \begin{cases} 0, & s \leq 0, \\ 1, & s > 0. \end{cases}$$



$$h_p(\mathbf{x}) = \text{step}\left(p(\mathbf{x}) - \frac{1}{2}\right)$$

How do we choose the probability function?

How do we build probability predictions?

Regression suggests using an intercept and one weight for each input feature.

$$f_{\boldsymbol{\theta}}(\mathbf{x}) = b + \mathbf{x}\mathbf{w} = b + \sum_{j=1}^d x_j w_j$$


$$b \in \mathbb{R}, \quad \mathbf{w} = (w_1, \dots, w_d)^{\top} \in \mathbb{R}^{d \times 1}$$

$$\boldsymbol{\theta} = \begin{pmatrix} b \\ \mathbf{w} \end{pmatrix} = (b, w_1, \dots, w_d)^{\top} \in \mathbb{R}^{(d+1) \times 1}$$

However, $f_{\boldsymbol{\theta}}(\mathbf{x})$ can be **any real number**!
Hence, it cannot serve directly as a probability.
Rather, $f_{\boldsymbol{\theta}}(\mathbf{x})$ is a **score** (also called a **logit**)
that we have to turn into a probability.

The sigmoid function

A smooth, increasing function turns each real score into a number between zero and one.



$$\sigma(s) := \frac{1}{1 + e^{-s}} = \frac{e^s}{1 + e^s}$$

$$0 < \sigma(s) < 1$$

$$\sigma(0) = \frac{1}{2}$$

Probabilities, labels and the boundary

Using the sigmoid σ we can turn scores into probabilities:

$$p_{\theta}(\mathbf{x}) := \sigma(f_{\theta}(\mathbf{x})) = \sigma(b + \mathbf{xw})$$

The corresponding **plug-in classifier** $h_{p_{\theta}}$ satisfies

$$h_{p_{\theta}}(\mathbf{x}) = 1 \quad \Leftrightarrow \quad p_{\theta}(\mathbf{x}) > \frac{1}{2} \quad \Leftrightarrow \quad f_{\theta}(\mathbf{x}) > 0 \quad \Leftrightarrow \quad b + \mathbf{xw} > 0$$

Decision boundary: the set of $\mathbf{x} \in \mathbb{R}^d$ such that $b + \mathbf{xw} = 0$

The **decision boundary** is a hyperplane with d features,
provided the weight vector $\mathbf{w}$ is nonzero.


If $d = 1$, $b + x_1 w_1 = 0$ determines a point of the Euclidean line

If $d = 2$, $b + x_1 w_1 + x_2 w_2 = 0$ determines a straight line of the Euclidean plane


If $d = 3$, $b + x_1 w_1 + x_2 w_2 + x_3 w_3 = 0$ determines a plane of the Euclidean space

Comparing two candidate boundaries

Which of these two classifiers is better?
Count the number of mistakes!



• Class 0 ▲ Class 1 — Decision boundary



• Class 0 ▲ Class 1 — Decision boundary

$$\theta = (-5.75, 0, 1)^T, \quad x_2 = 5.75 \quad \theta = (-1, -0.6, 1)^T, \quad x_2 = 0.6x_1 + 1$$

The classifier on the right makes fewer mistakes on these observations

Today's lecture

Binary classification

Classification error and logistic loss

Gradient descent

More than two classes

Finding the parameters

We want the **smallest mistake count**.

Finding the parameter θ^* that minimizes the empirical 0-1 loss:

$$\theta \mapsto \mathcal{E}(\theta) = \sum_{i=1}^N \ell_{0,1}(h_{p_\theta}(\mathbf{x}_i), y_i)$$

is a **combinatorial** problem.



Can we compute an **approximate solution** via **gradient descent** as in regression?

$$\theta^{(t+1)} = \theta^{(t)} - \eta \nabla \mathcal{E}(\theta^{(t)}), \quad \eta > 0$$

What happens to the mistake count when we move the boundary a little?

Moving the boundary and counting mistakes

Fix the weights and vary only the intercept. The line moves parallel to itself.



$$b = -5.55, \quad w_1 = 0, \quad w_2 = 1$$

The count is **locally constant** and changes only when the line crosses an observed point.

Flat regions give no direction

If every observed score is **nonzero**
small parameter changes preserve **the signs** of the scores
and hence **all predicted labels stay the same**.

$$\mathcal{E}(\boldsymbol{\theta} + \Delta) = \mathcal{E}(\boldsymbol{\theta}) \quad \text{for sufficiently small } \Delta$$

The **gradient** of a **locally constant** function is **zero**!

$$\nabla \mathcal{E}(\boldsymbol{\theta}) = \mathbf{0}$$

Hence

$$\boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \eta \nabla \mathcal{E}(\boldsymbol{\theta}^{(t)}) = \boldsymbol{\theta}^{(t)}$$

Gradient descent will **not move!!!**

What does the mistake count miss?

Suppose the **observed class is 0**. Both predictions below choose class 1.

	Prediction A	Prediction B
Probability of class 1	0.51	0.90
Predicted class	1	1
Zero-one loss	1	1
Probability of the observed class	0.49	0.10

At 0.51, the model is almost guessing: we want a milder penalty.

At 0.90, it is confidently wrong: we want a larger penalty.

Rather than scoring only the predicted label,
we **penalise assigning low probability to the observed class**.

Logistic loss: a surrogate for fitting probabilities

For $0 < p < 1$ and $y \in \{0, 1\}$, take the **negative logarithm** of the probability assigned to the observed class.

$$\ell_{\log}(p, y) = -y\log(p) - (1 - y)\log(1 - p) = \begin{cases} -\log(p), & y = 1, \\ -\log(1 - p), & y = 0. \end{cases}$$

The **logistic loss** makes us pay the negative logarithm of the probability we assign to the correct label.

Wrong, but uncertain

$$\begin{aligned} \ell_{\log}(0.51, 0) &= -\log(0.49) \\ &\approx 0.713 \end{aligned}$$

Confidently wrong

$$\begin{aligned} \ell_{\log}(0.90, 0) &= -\log(0.10) \\ &\approx 2.303 \end{aligned}$$

Logistic loss is also called **binary cross-entropy**. We use natural logarithms.

The empirical logistic loss

The observations stay fixed.

The **probability predictions** depend on the parameters.

$$p_{\theta}(\mathbf{x}_i) = \sigma(b + \mathbf{x}_i \mathbf{w})$$


$$\mathcal{L}_{\log}(\theta) = \sum_{i=1}^N \ell_{\log}(p_{\theta}(\mathbf{x}_i), y_i)$$

We fit this **smooth surrogate loss using **gradient descent**, while the mistake count $\varepsilon(\theta)$ remains our classification measure.**

Comparing the two losses


A mistake gives the observed class probability at most $1/2$,
so its logistic loss is at least $-\log\left(\frac{1}{2}\right) = \log(2)$.

Observed label $y = 0$



— Logistic loss / log 2 — Zero-one loss

Observed label $y = 1$



— Logistic loss / log 2 — Zero-one loss

$$\ell_{0,1}\left(\text{step}\left(\sigma(s) - \frac{1}{2}\right), y\right) \leq \frac{\ell_{\log}(\sigma(s), y)}{\log(2)}$$

Bounding the number of mistakes

The inequality holds for each observation

We can therefore add it over the dataset

$$\begin{aligned}\mathcal{E}(\boldsymbol{\theta}) &= \sum_{i=1}^N \ell_{0,1}(h_{p_{\boldsymbol{\theta}}}(\mathbf{x}_i), y_i) \\ &= \sum_{i=1}^N \ell_{0,1}(\text{step}(f_{\boldsymbol{\theta}}(\mathbf{x}_i)), y_i) \leq \sum_{i=1}^N \frac{\ell_{\log}(\sigma(f_{\boldsymbol{\theta}}(\mathbf{x}_i)), y_i)}{\log(2)} \\ &= \frac{1}{\log(2)} \sum_{i=1}^N \ell_{\log}(p_{\boldsymbol{\theta}}(\mathbf{x}_i), y_i) = \frac{1}{\log(2)} \cdot \mathcal{L}_{\log}(\boldsymbol{\theta})\end{aligned}$$

A control over the empirical logistic loss
gives a control on the number of mistakes

Today's lecture

Binary classification

Classification error and logistic loss

Gradient descent

More than two classes

The design matrix and the parameters

As in Lecture 2, a first column of ones includes the intercept.

$$\mathbf{X} = \begin{pmatrix} 1 & x_{11} & \cdots & x_{1d} \\ \vdots & \vdots & & \vdots \\ 1 & x_{N1} & \cdots & x_{Nd} \end{pmatrix} \quad \boldsymbol{\theta} = \begin{pmatrix} b \\ w_1 \\ \vdots \\ w_d \end{pmatrix}$$

$$\mathbf{X} \in \mathbb{R}^{N \times (d+1)}, \quad \boldsymbol{\theta} \in \mathbb{R}^{(d+1) \times 1}$$

$$\theta_1 = b, \quad \theta_{k+1} = w_k, \quad k = 1, \dots, d$$

The first coordinate is the intercept. The remaining coordinates are the weights.

The three functions

Fix the i th input/output pair $(\mathbf{x}_i, y_i)$ and vary the model parameters $\boldsymbol{\theta}$.

Define affine score corresponding to the i th input $\mathbf{x}_i$

$$F_i(\boldsymbol{\theta}) := f_{\boldsymbol{\theta}}(\mathbf{x}_i) = b + \mathbf{x}_i \mathbf{w} = (1, \mathbf{x}_i) \boldsymbol{\theta} = \sum_{r=1}^{d+1} X_{ir} \theta_r$$

Recall the definition of the sigmoid function σ

$$\sigma(s) = \frac{1}{1 + e^{-s}} \quad (s \in \mathbb{R})$$

Define the logistic loss corresponding to the i th output y_i

$$G_i(p) := \ell_{\log}(p, y_i)$$

The chain rule for one observation

The loss for the i th observation is the **composition** of these three functions.

$$\ell_{\log}(p_{\boldsymbol{\theta}}(\mathbf{x}_i), y_i) = (G_i \circ \sigma \circ F_i)(\boldsymbol{\theta}) = G_i \left(\sigma(F_i(\boldsymbol{\theta})) \right)$$

To compute the partial derivatives with respect to θ_j we apply the **chain rule to this composition.**

$$\begin{aligned} \frac{\partial}{\partial \theta_j} \ell_{\log}(p_{\boldsymbol{\theta}}(\mathbf{x}_i), y_i) &= \frac{\partial}{\partial \theta_j} (G_i \circ \sigma \circ F_i)(\boldsymbol{\theta}) = G'_i \left(\sigma(F_i(\boldsymbol{\theta})) \right) \cdot \sigma'(F_i(\boldsymbol{\theta})) \cdot \frac{\partial F_i}{\partial \theta_j}(\boldsymbol{\theta}) \\ &= \frac{\partial \ell_{\log}}{\partial p} \left(\sigma(f_{\boldsymbol{\theta}}(\mathbf{x}_i)), y_i \right) \cdot \sigma'(f_{\boldsymbol{\theta}}(\mathbf{x}_i)) \cdot \frac{\partial f_{\boldsymbol{\theta}}(\mathbf{x}_i)}{\partial \theta_j} \\ &= \underbrace{\frac{\partial \ell_{\log}}{\partial p} (p_{\boldsymbol{\theta}}(\mathbf{x}_i), y_i)}_{\text{first term}} \cdot \underbrace{\sigma'(f_{\boldsymbol{\theta}}(\mathbf{x}_i))}_{\text{second term}} \cdot \underbrace{\frac{\partial f_{\boldsymbol{\theta}}(\mathbf{x}_i)}{\partial \theta_j}}_{\text{third term}} \end{aligned}$$

First term: Differentiating the logistic loss

Hold the observed label fixed and vary the probability.

$$\ell_{\log}(p, y) = -y\log(p) - (1 - y)\log(1 - p)$$

$$\begin{aligned}\frac{\partial \ell_{\log}}{\partial p}(p, y) &= -\frac{y}{p} + \frac{1 - y}{1 - p} = \frac{-y(1 - p) + (1 - y)p}{p(1 - p)} \\ &= \frac{-y + yp + p - yp}{p(1 - p)} = \frac{p - y}{p(1 - p)}\end{aligned}$$

$$\boxed{\frac{\partial \ell_{\log}}{\partial p}(p_{\theta}(\mathbf{x}_i), y_i) = \frac{p_{\theta}(\mathbf{x}_i) - y_i}{p_{\theta}(\mathbf{x}_i)(1 - p_{\theta}(\mathbf{x}_i))}}$$

Second term: Differentiating the sigmoid

Write the **sigmoid as a composition** and apply the **chain rule** again.

Inner function u

$$u(s) = 1 + e^{-s}, \quad u'(s) = -e^{-s}$$

Outer function v

$$v(z) = \frac{1}{z}, \quad v'(z) = -\frac{1}{z^2}$$

$$\sigma(s) = \frac{1}{1 + e^{-s}} = v(u(s))$$

$$\sigma'(s) = v'(u(s)) u'(s) = -\frac{1}{(1 + e^{-s})^2} (-e^{-s}) = \frac{e^{-s}}{(1 + e^{-s})^2}$$

$$= \frac{1}{1 + e^{-s}} \cdot \frac{e^{-s}}{1 + e^{-s}} = \frac{1}{1 + e^{-s}} \cdot \left(1 - \frac{1}{1 + e^{-s}}\right) = \sigma(s)(1 - \sigma(s))$$

$$\boxed{\sigma'(f_{\theta}(\mathbf{x}_i)) = \sigma(f_{\theta}(\mathbf{x}_i)) (1 - \sigma(f_{\theta}(\mathbf{x}_i))) = p_{\theta}(\mathbf{x}_i)(1 - p_{\theta}(\mathbf{x}_i))}$$

Third term: Differentiating the affine score

The inputs are fixed, so only the matching parameter contributes.

The affine score

$$f_{\boldsymbol{\theta}}(\mathbf{x}_i) = \sum_{r=1}^{d+1} X_{ir} \theta_r$$

The Kronecker delta

$$\frac{\partial \theta_r}{\partial \theta_j} = \delta_{rj} = \begin{cases} 1, & r = j, \\ 0, & r \neq j. \end{cases}$$

$$\frac{\partial f_{\boldsymbol{\theta}}(\mathbf{x}_i)}{\partial \theta_j} = \frac{\partial}{\partial \theta_j} \sum_{r=1}^{d+1} X_{ir} \theta_r = \sum_{r=1}^{d+1} X_{ir} \frac{\partial \theta_r}{\partial \theta_j} = \sum_{r=1}^{d+1} X_{ir} \delta_{rj} = X_{ij}$$

$$\boxed{\frac{\partial f_{\boldsymbol{\theta}}(\mathbf{x}_i)}{\partial \theta_j} = X_{ij}} \quad j = 1, \dots, d + 1$$

Putting the three terms together

Substitute the first, second and third terms into the chain rule.

$$\begin{aligned} & \frac{\partial}{\partial \theta_j} \ell_{\log}(p_{\boldsymbol{\theta}}(\mathbf{x}_i), y_i) \\ &= \underbrace{\frac{\partial \ell_{\log}}{\partial p}(p_{\boldsymbol{\theta}}(\mathbf{x}_i), y_i)}_{\text{first term}} \cdot \underbrace{\sigma'(f_{\boldsymbol{\theta}}(\mathbf{x}_i))}_{\text{second term}} \cdot \underbrace{\frac{\partial f_{\boldsymbol{\theta}}(\mathbf{x}_i)}{\partial \theta_j}}_{\text{third term}} \\ &= \underbrace{\frac{p_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i}{p_{\boldsymbol{\theta}}(\mathbf{x}_i)(1 - p_{\boldsymbol{\theta}}(\mathbf{x}_i))}}_{\text{first term}} \cdot \underbrace{p_{\boldsymbol{\theta}}(\mathbf{x}_i)(1 - p_{\boldsymbol{\theta}}(\mathbf{x}_i))}_{\text{second term}} \cdot \underbrace{X_{ij}}_{\text{third term}} \\ &= (p_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i)X_{ij} \end{aligned}$$

Empirical gradient and gradient descent

The empirical loss is a sum, so we add the observation derivatives.

$$\frac{\partial \mathcal{L}_{\log}}{\partial \theta_j}(\boldsymbol{\theta}) = \sum_{i=1}^N \frac{\partial}{\partial \theta_j} \ell_{\log}(p_{\boldsymbol{\theta}}(\mathbf{x}_i), y_i) = \sum_{i=1}^N X_{ij}(p_{\boldsymbol{\theta}}(\mathbf{x}_i) - y_i)$$

Sigmoid applied to each entry

$$\boldsymbol{\sigma}((s_1, \dots, s_N)^\top) := (\sigma(s_1), \dots, \sigma(s_N))^\top$$

One probability per observation

$$\mathbf{p}(\boldsymbol{\theta}) := (p_{\boldsymbol{\theta}}(\mathbf{x}_1), \dots, p_{\boldsymbol{\theta}}(\mathbf{x}_N))^\top$$

$$\mathbf{p}(\boldsymbol{\theta}) = \boldsymbol{\sigma}(\mathbf{X}\boldsymbol{\theta}), \quad \mathbf{y} = (y_1, \dots, y_N)^\top$$

Gradient

$$\nabla \mathcal{L}_{\log}(\boldsymbol{\theta}) = \mathbf{X}^\top(\mathbf{p}(\boldsymbol{\theta}) - \mathbf{y}) = \mathbf{X}^\top(\boldsymbol{\sigma}(\mathbf{X}\boldsymbol{\theta}) - \mathbf{y})$$

Fitting logistic regression

$$\boldsymbol{\theta}^{(t+1)} = \boldsymbol{\theta}^{(t)} - \eta \nabla \mathcal{L}_{\log}(\boldsymbol{\theta}^{(t)}) = \boldsymbol{\theta}^{(t)} - \eta \mathbf{X}^\top(\boldsymbol{\sigma}(\mathbf{X}\boldsymbol{\theta}^{(t)}) - \mathbf{y})$$

Stationary points and minimisers

Fact: the empirical logistic loss is **convex** in the parameters.
Therefore, **any stationary point is a global minimiser.**

However: a finite minimiser need not exist!!!
See the notes for a counterexample.


Warning 1: decreasing the empirical logistic loss
does not guarantee fewer mistakes (a lower empirical 0–1 loss).

Warning 2: gradient descent still needs an appropriate learning rate.
A step that is too large can increase the loss.

Gradient descent: $\eta = 0.001, t = 0$


Initial coefficients: $b^{(0)} = -5.75,$ $w_1^{(0)} = 0,$ $w_2^{(0)} = 1$

Learning rate: $\eta = 0.001$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 5.7500 + 0.0000x_1$$




—●— Logistic loss / log 2 —●— Number of mistakes


Logistic loss: 51.048
Mistakes: 31 of 100

Gradient descent: $\eta = 0.001, t = 1$

Coefficients: $b^{(1)} \approx -5.732582$ $w_1^{(1)} \approx 0.116619$ $w_2^{(1)} \approx 1.105206$



● Class 0 ▲ Class 1 — Decision boundary




— Logistic loss / log 2 — Number of mistakes

$$x_2 = 5.1869 - 0.1055x_1$$

Logistic loss: 42.733
Mistakes: 15 of 100


Gradient descent: $\eta = 0.001, t = 2$

Coefficients: $b^{(2)} \approx -5.740878$ $w_1^{(2)} \approx 0.070300$ $w_2^{(2)} \approx 1.081132$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 5.3101 - 0.0650x_1$$




— Logistic loss / log 2 — Number of mistakes

Logistic loss: 41.668
Mistakes: 7 of 100


Gradient descent: $\eta = 0.001, t = 5$

Coefficients: $b^{(5)} \approx -5.744795$ $w_1^{(5)} \approx 0.062560$ $w_2^{(5)} \approx 1.111266$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 5.1696 - 0.0563x_1$$




— Logistic loss / log 2 — Number of mistakes

Logistic loss: 41.300
Mistakes: 8 of 100


Gradient descent: $\eta = 0.001, t = 10$

Coefficients: $b^{(10)} \approx -5.753743$ $w_1^{(10)} \approx 0.034272$ $w_2^{(10)} \approx 1.148470$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 5.0099 - 0.0298x_1$$




— Logistic loss / log 2 — Number of mistakes

Logistic loss: 40.848
Mistakes: 8 of 100


Gradient descent: $\eta = 0.001, t = 20$

Coefficients: $b^{(20)} \approx -5.771210$ $w_1^{(20)} \approx -0.020164$ $w_2^{(20)} \approx 1.221302$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 4.7255 + 0.0165x_1$$




—●— Logistic loss / log 2 —●— Number of mistakes

Logistic loss: 39.992
Mistakes: 7 of 100

The loss can fall while mistakes increase

The same planar run shows the limitation of the surrogate.



Step	Logistic loss	Errors
3	41.4948	7
4	41.3932	8


Of course: the bound still holds at both parameter choices.

Logistic loss can decrease even when the mistake count increases.


Gradient descent: $\eta = 0.002, t = 0$

Initial coefficients: $b^{(0)} = -5.75,$ $w_1^{(0)} = 0,$ $w_2^{(0)} = 1$

Learning rate: $\eta = 0.002$



● Class 0 ▲ Class 1 — Decision boundary




— Logistic loss / log 2 — Number of mistakes

$$x_2 = 5.7500 + 0.0000x_1$$

Logistic loss: 51.048
Mistakes: 31 of 100


Gradient descent: $\eta = 0.002, t = 1$

Coefficients: $b^{(1)} \approx -5.715164$ $w_1^{(1)} \approx 0.233238$ $w_2^{(1)} \approx 1.210411$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 4.7217 - 0.1927x_1$$




— Logistic loss / log 2 — Number of mistakes

Logistic loss: 64.556
Mistakes: 42 of 100


Gradient descent: $\eta = 0.002, t = 2$

Coefficients: $b^{(2)} \approx -5.773416$ $w_1^{(2)} \approx -0.115049$ $w_2^{(2)} \approx 0.966988$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 5.9705 + 0.1190x_1$$




—●— Logistic loss / log 2 —●— Number of mistakes

Logistic loss: 75.928
Mistakes: 49 of 100


Gradient descent: $\eta = 0.002, t = 5$

Coefficients: $b^{(5)} \approx -5.718821$ $w_1^{(5)} \approx 0.295982$ $w_2^{(5)} \approx 1.392324$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 4.1074 - 0.2126x_1$$




— Logistic loss / log 2 — Number of mistakes

Logistic loss: 105.343
Mistakes: 49 of 100


Gradient descent: $\eta = 0.002, t = 10$

Coefficients: $b^{(10)} \approx -5.825631$ $w_1^{(10)} \approx -0.256765$ $w_2^{(10)} \approx 1.132075$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 5.1460 + 0.2268x_1$$




— Logistic loss / log 2 — Number of mistakes

Logistic loss: 77.632
Mistakes: 49 of 100


Gradient descent: $\eta = 0.002, t = 20$

Coefficients: $b^{(20)} \approx -5.884803$ $w_1^{(20)} \approx -0.405317$ $w_2^{(20)} \approx 1.344522$



● Class 0 ▲ Class 1 — Decision boundary

$$x_2 = 4.3769 + 0.3015x_1$$




— Logistic loss / log 2 — Number of mistakes

Logistic loss: 72.282
Mistakes: 48 of 100

Step size and logistic loss

Same data and initial coefficients, only the learning rate changes.



Step	Loss $\eta = 0.001$	Loss $\eta = 0.002$
0	51.048	51.048
1	42.733	64.556

**At $\eta = 0.002$,
even the first update
increases the loss.**

A step that is too large can increase the loss.

Today's lecture

Binary classification


Classification error and logistic loss

Gradient descent

More than two classes

More than two classes

Recognising a handwritten digit means choosing among ten classes: 0, 1, ..., 9.



Each row contains examples of the same digit.

An image is made of pixels.
In greyscale, their intensities
form a matrix of real numbers.

We read the matrix row by row
to form a row vector.

We use the same pixel order
for every image.

$$\mathbf{x} = (x_1, \dots, x_d) \in \mathbb{R}^{1 \times d}$$

Here, d is the number of pixels.

We predict one probability per class

Predicting class probabilities

We number the classes from 1 to K .

For digits 0 to 9, class index k refers to digit $k - 1$.

$$\mathbf{p}: \mathbb{R}^{1 \times d} \rightarrow [0,1]^{1 \times K}$$

$$\mathbf{p}(\mathbf{x}) = (p_1(\mathbf{x}), \dots, p_K(\mathbf{x}))$$

Each component gives the predicted probability of the corresponding class.

For this vector to represent a valid probability distribution,
we require its components to sum to one:

$$\sum_{k=1}^K p_k(\mathbf{x}) = 1$$

If we must choose one label, which should we predict?

The multiclass plug-in rule

Choose the class with the largest predicted probability.

$$h_{\mathbf{p}}: \mathbb{R}^{1 \times d} \longrightarrow \{1, \dots, K\}$$

$$p_{h_{\mathbf{p}}(\mathbf{x})}(\mathbf{x}) = \max_{k \in \{1, \dots, K\}} p_k(\mathbf{x})$$

We are going to return the smallest index in case of a tie.

The largest probability need not exceed one half.

How do we build the probability function?

One affine score per class

for each class $k = 1, \dots, K$

$$b_k \in \mathbb{R}, \quad \mathbf{w}_k = \begin{pmatrix} w_{1k} \\ \vdots \\ w_{dk} \end{pmatrix} \in \mathbb{R}^{d \times 1} \quad \boldsymbol{\theta}_k = \begin{pmatrix} b_k \\ \mathbf{w}_k \end{pmatrix} \in \mathbb{R}^{(d+1) \times 1}$$

The k -th score function is defined via

$$f_{\boldsymbol{\theta}_k}(\mathbf{x}) = b_k + \mathbf{x}\mathbf{w}_k$$

As in binary classification, these affine scores are also called **logits**.

Stack all the parameter columns in a matrix Θ

$$\Theta = (\boldsymbol{\theta}_1 \quad \dots \quad \boldsymbol{\theta}_K) \in \mathbb{R}^{(d+1) \times K}$$

$$\mathbf{f}_{\Theta}(\mathbf{x}) = \left(f_{\boldsymbol{\theta}_1}(\mathbf{x}), \dots, f_{\boldsymbol{\theta}_K}(\mathbf{x}) \right) \in \mathbb{R}^{1 \times K}$$

Softmax turns scores into probabilities

Warning: separate sigmoids do not ensure a sum of one!!!

Three zero scores would give three probabilities of one half.

Exponentiate each score, then divide by the sum across classes.

$$\text{softmax}: \mathbb{R}^{1 \times K} \rightarrow (0,1)^{1 \times K}$$

$$\mathbf{s} = (s_1, \dots, s_K) \in \mathbb{R}^{1 \times K}$$

$$[\text{softmax}(\mathbf{s})]_k = \frac{e^{s_k}}{\sum_{r=1}^K e^{s_r}}, \quad k = 1, \dots, K$$

Positive numbers:

$$[\text{softmax}(\mathbf{s})]_k > 0$$

Valid probabilities:

$$\sum_{k=1}^K [\text{softmax}(\mathbf{s})]_k = \frac{\sum_{k=1}^K e^{s_k}}{\sum_{r=1}^K e^{s_r}} = 1$$

Our multiclass probability model

Apply softmax to the row of affine scores.

$$\mathbf{p}_{\Theta}(\mathbf{x}) = \text{softmax}(\mathbf{f}_{\Theta}(\mathbf{x}))$$

Its component for each class is

$$p_{\Theta,k}(\mathbf{x}) := [\mathbf{p}_{\Theta}(\mathbf{x})]_k = \frac{\exp(f_{\Theta_k}(\mathbf{x}))}{\sum_{r=1}^K \exp(f_{\Theta_r}(\mathbf{x}))}, \quad k = 1, \dots, K$$

Once we have the probability predictor $\mathbf{p}_{\Theta}$

We can build the plug-in classifier $h_{\mathbf{p}_{\Theta}}$

Cross-entropy loss for the observed class

The loss is minus the logarithm of the predicted probability of the observed class

$$\mathbf{p} = (p_1, \dots, p_K), \quad y \in \{1, \dots, K\}, \quad p_y > 0$$

$$\ell_{\text{multi}}(\mathbf{p}, y) = -\log(p_y)$$

Dataset: $(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_N, y_N) \in \mathbb{R}^{1 \times d} \times \{1, \dots, K\}$

Empirical cross-entropy

$$\mathcal{L}_{\text{multi}}(\Theta) = \sum_{i=1}^N \ell_{\text{multi}}(\mathbf{p}_{\Theta}(\mathbf{x}_i), y_i)$$

**Fit the parameter matrix by minimising this summed loss,
e.g., via **gradient descent**.**

Models and fitting losses

Task	Model predictions	Fitting loss
Regression	Numerical values	Squared loss
Binary classification	Affine/Sigmoid probabilities	Logistic loss
Multiclass classification	Affine/Softmax probabilities	Cross-entropy loss

We fit probabilities with a smooth surrogate loss.
We evaluate classification by counting mistakes.

Lecture notes and exercises explain the material in more detail.
Exercise solutions will be posted before the next lecture.

github.com/semtm0051/2026 27

